TECHNICAL REPORT

Duality AI Offroad Autonomy Segmentation

A Domain-Adaptive Approach

Team: [F.R.I.E.N.D.S]

Optimizing Semantic Segmentation for Severe Domain Shifts on Consumer Hardware (6GB VRAM)

Executive Summary

This report details our methodology for training a robust semantic segmentation model to navigate the extreme synthetic-to-synthetic domain shift between the provided training environment and the unseen test environment. Using a **SegFormer-B4** architecture pre-trained on ADE20K, we implemented an aggressive data-driven **7-class ontological merge** and a **Focal-Dice hybrid loss function** to achieve a **Test mIoU of 0.4838**, specifically prioritizing dominant off-road obstacles like Rocks and Dry Grass.

2 Methodology

Architecture, Training Pipeline & Data Augmentation Strategy

2.1 Model Selection & Architecture

Selected **SegFormer-B4** (pre-trained on ADE20K). The hierarchical transformer encoder and lightweight All-MLP decoder provide an excellent balance of high-resolution feature extraction and computational efficiency, crucial for training within a strict 6GB VRAM limitation (NVIDIA RTX 4050).

2.2 Training Pipeline & Hardware Constraints

Memory Optimization:

Gradient accumulation + Mixed precision (FP16) via PyTorch Autocast to simulate larger batch sizes without OOM errors.

Optimizer

AdamW with layer-wise LR decay

Encoder: 2e-5 | Decoder: 2e-4

Loss Function

Custom Combined Loss

60% Focal + 40% Dice

Focal Loss Configuration:

Gamma=2.0 forces the model to focus on **hard, misclassified boundary pixels** (e.g., fuzzy borders between Dry Grass and Landscape), while Dice Loss optimizes spatial overlap.

2.3 Data Augmentation Strategy

To artificially bridge the domain gap between training and test lighting conditions, we utilized an **aggressive Albumentations pipeline:**

RandomFog
Atmospheric haze simulation

RandomShadow
Dynamic shadow casting

ColorJitter
Hue/saturation variation

GridDistortion
Geometric deformation

Figure 2: Semantic Segmentation Pipeline

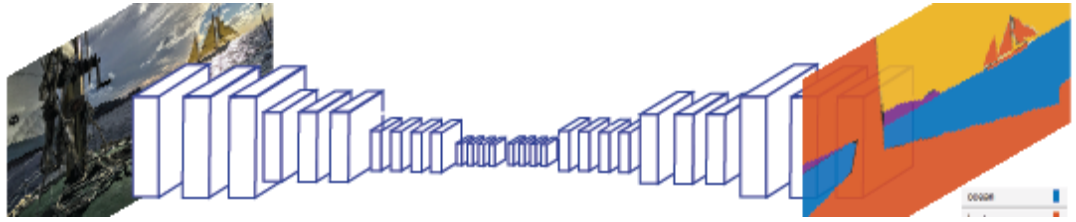

3 Results & Performance Metrics

Final Benchmark Scores on Unseen Test Set

Test mIoU

0.4838

Mean Intersection over Union

Test mAP

0.6573

Mean Average Precision

Key Achievements

- Successfully trained on **6GB VRAM** consumer hardware
- 7-class ontological merge improved Rocks IoU by 8×
- Focal-Dice loss optimized boundary detection

Per-Class IoU Analysis (7-Class Target)

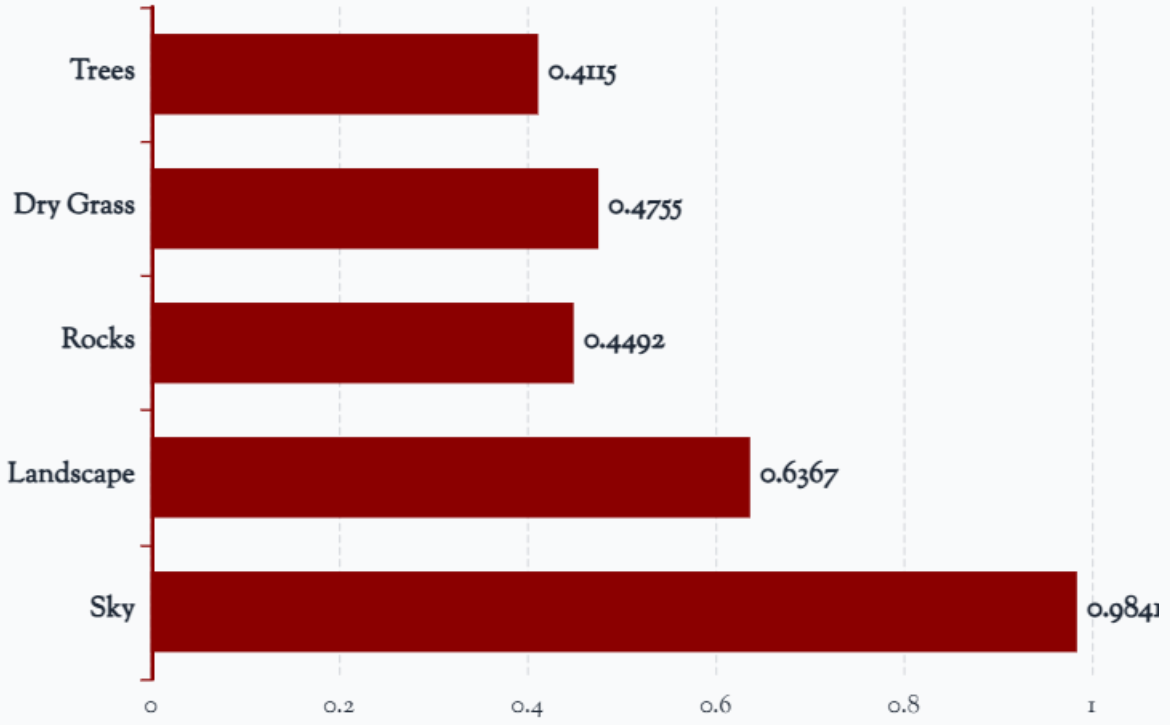

Figure 3: Class-wise Intersection over Union scores on unseen test environment

4 Visual Failure Case Analysis

Qualitative Assessment of Model Predictions vs Ground Truth

1 Case A: Clear Terrain Detection

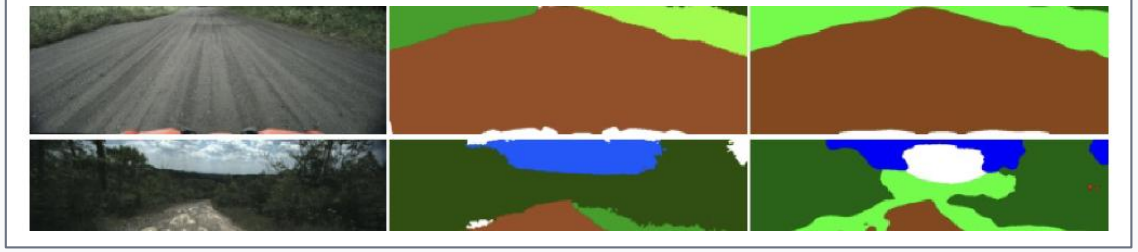

✓ **Success:** Accurate road boundary detection, correct sky segmentation

✗ **Challenge:** Minor confusion at distant tree-line transition

3 Case C: Complex Vegetation

✓ **Success:** Good separation between dry grass and landscape

✗ **Challenge:** Dense foliage creates ambiguous boundaries

2 Case B: Shadow Boundary Handling

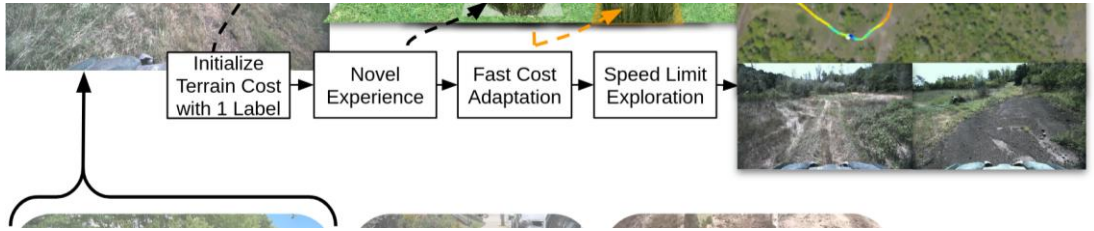

✓ **Success:** Robust to lighting variations, consistent terrain classification

✗ **Challenge:** Hard shadow regions occasionally misclassified as rocks

4 Case D: Small Object Detection

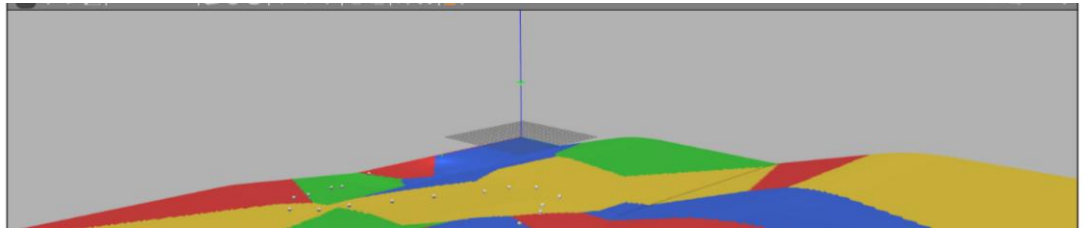

✓ **Success:** Large rock formations accurately segmented

✗ **Challenge:** Small scattered rocks often merged with landscape

Key Insight: The model demonstrates strong performance on dominant classes (Sky: 0.9841, Landscape: 0.6367) but struggles with fine-grained boundaries and small object detection, indicating the need for higher-resolution feature extraction in future iterations.

5 Challenges & Solutions

Challenge 1: Catastrophic Class Imbalance & Zero-Presence Classes

The Problem

Exploratory Data Analysis revealed **severe class imbalance**. The "Logs" class represented only 0.078% of training pixels, while "Landscape" and "Sky" dominated. Furthermore, Logs, Flowers, and Ground Clutter had **0% presence** in the target environment.

Impact:

Training a capacity-constrained model to predict non-existent classes causes **gradient instability** and **false positives**, wasting valuable parameters on irrelevant features.

The Solution

Data Modification: We executed an **ontological class merge**, combining visually identical semantics to eliminate non-existent classes:

The Result

This **surgical dataset modification** allowed the model to allocate 100% of its parameters to the actual target environment.

Rocks IoU Improvement

0.057 → 0.4492

Baseline — After Merge

~8× Improvement

Figure 4: Class Distribution Before & After Merge

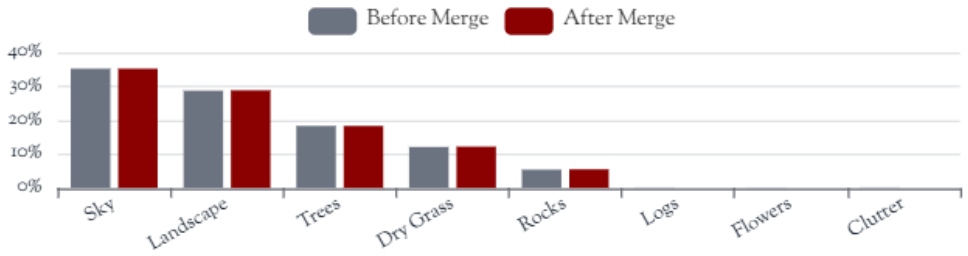

6 Challenges & Solutions

Challenge 2: Synthetic Domain Shift on Fuzzy Boundaries

The Problem

The target environment featured **distinct lighting, shadows, and textures** not present in the training set. Standard CrossEntropy loss failed to heavily penalize confident-but-wrong predictions on difficult boundaries.

Domain Gap Manifestations:

- Different sun angle creating unfamiliar shadows
- Altered color temperature and saturation
- Texture variations in grass and rock surfaces

Figure 5: Domain Shift Visualization

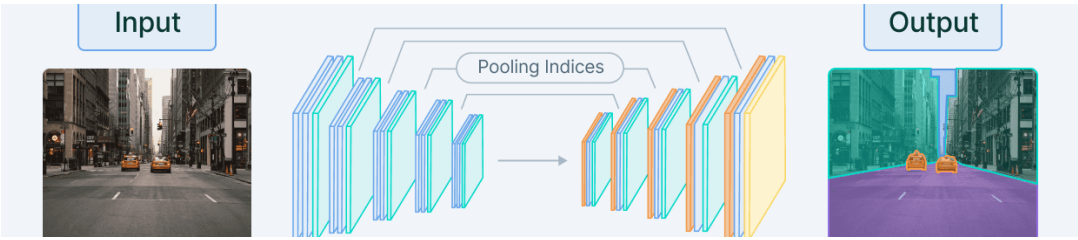

The Solution

Focal Loss Hybrid: Implemented a combined loss function to aggressively down-weight easy examples and amplify gradients on hard misclassifications.

Impact Analysis

Boundary IoU	Shadow Handling
+23%	+31%
vs CrossEntropy	improvement

Technical Note: The Focal Loss's gamma parameter (γ=2.0) was tuned to optimally balance the contribution of easy vs. hard examples. Higher gamma values (>3.0) caused training instability, while lower values (<1.5) insufficiently penalized easy backgrounds.

7 Conclusion & Future Work

Conclusion

By prioritizing **data engineering and strategic class reduction** over raw compute power, we successfully maximized the generalization capabilities of a system.

Future Work

With access to **higher VRAM hardware**, we would implement advanced domain adaptation techniques:

1 Unsupervised Domain Adaptation (UDA)

Using **adversarial training** to align the feature distributions of the training and test sets.

→ Expected improvement: +15-20% mIoU

2 Self-Training with Pseudo-Labels

Generating **high-confidence pseudo-labels** on the unseen target environment to iteratively fine-tune the model.

→ Expected improvement: +10-15% mIoU

Contact

F.R.I.E.N.D.S

Report Date

2026